



# Blocking vs non-blocking I/O




# Blocking vs non-blocking I/O



# Blocking vs non-blocking I/O

## blocking socket


thread is suspended until read/write from/to the socket completes



## blocking I/O

## non-blocking socket

thread reads data available in the socket buffer and does not wait for the remaining data to arrive

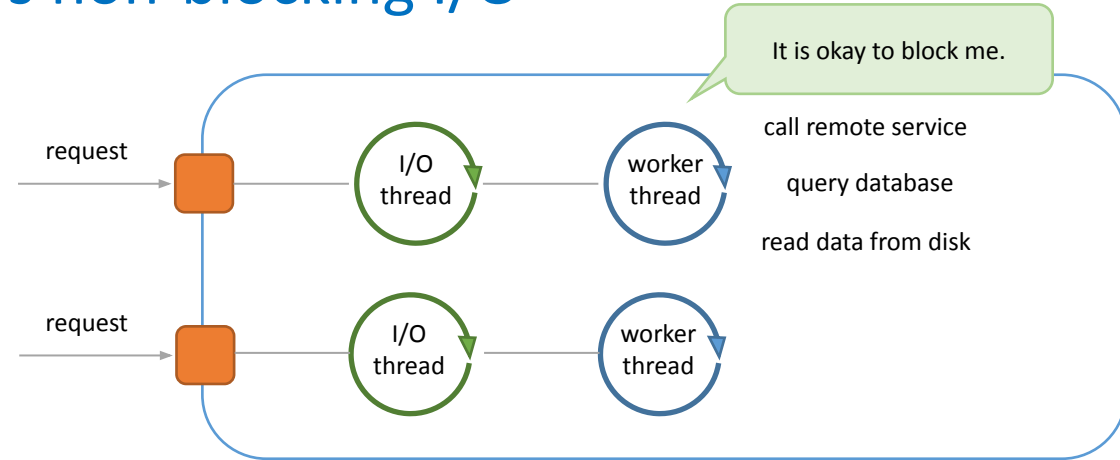


## non-blocking I/O

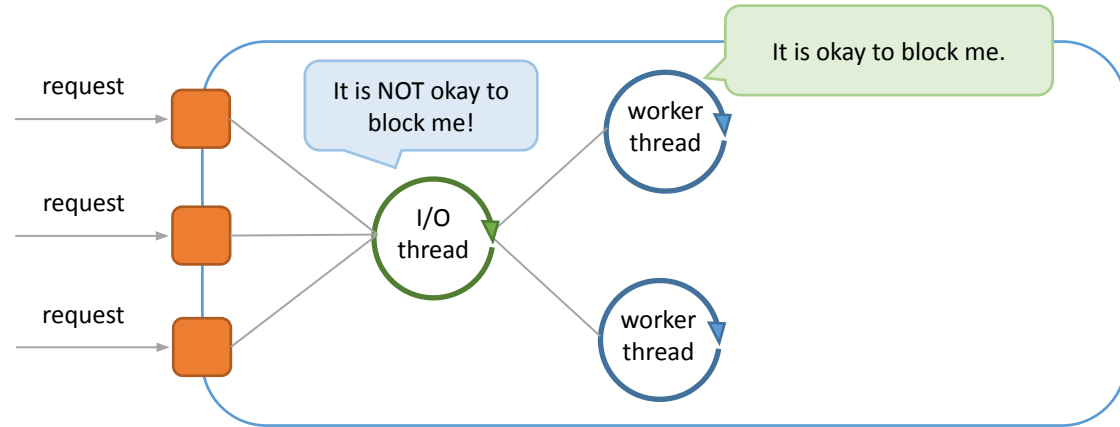
allows the server to handle a lot more connections than blocking I/O, since connections are cheap and threads are not

# Blocking vs non-blocking I/O

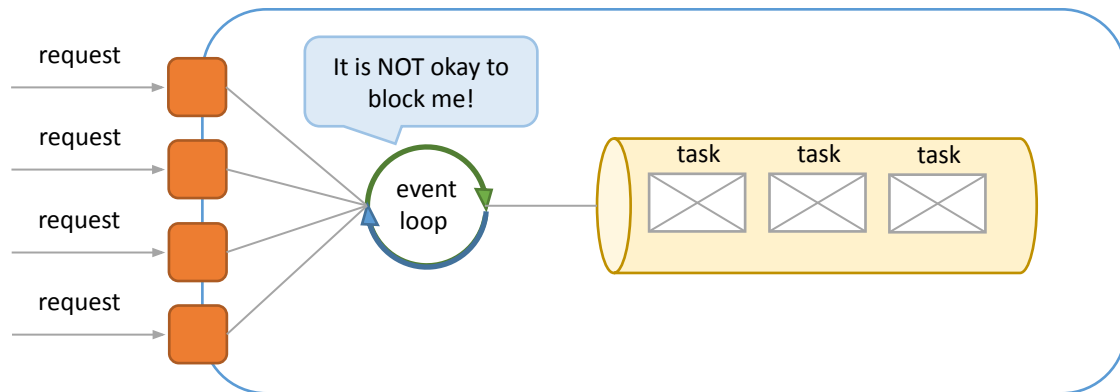
thread per connection



thread per request  
with non-blocking I/O




event loop




# Blocking vs non-blocking I/O

## concurrency



## parallelism



the optimal number of event loops is  
the number of CPUs

# Blocking vs non-blocking I/O

outdated  
still used in older versions of HTTP servers

## thread per connection



Jetty, Nginx, Tomcat

## thread per request

with non-blocking I/O

- Easier to implement, test and debug applications.
- Each thread consumes resources (memory, CPU).  
We have to limit the number of concurrent requests  
(load shedding and rate limiting).

good for  
CPU-bound (compute-intensive) workloads



Netty, Node.js, Zuul

## event loop

- Massive amount of active connections.
- More resilient to sudden traffic spikes.
- Increased development and operational complexity.

good for  
I/O-bound workloads (e.g. large files)