


# Log

## pros

- Put and get message operations are fast.



## cons


- **Limited memory space.**  
becomes a problem when
  - consumers get slower
  - messages are big
  - producers start to send more messages
- **Low durability guarantee.**

# Log

in memory

queue data structure

add



...



on disk

file

append



...

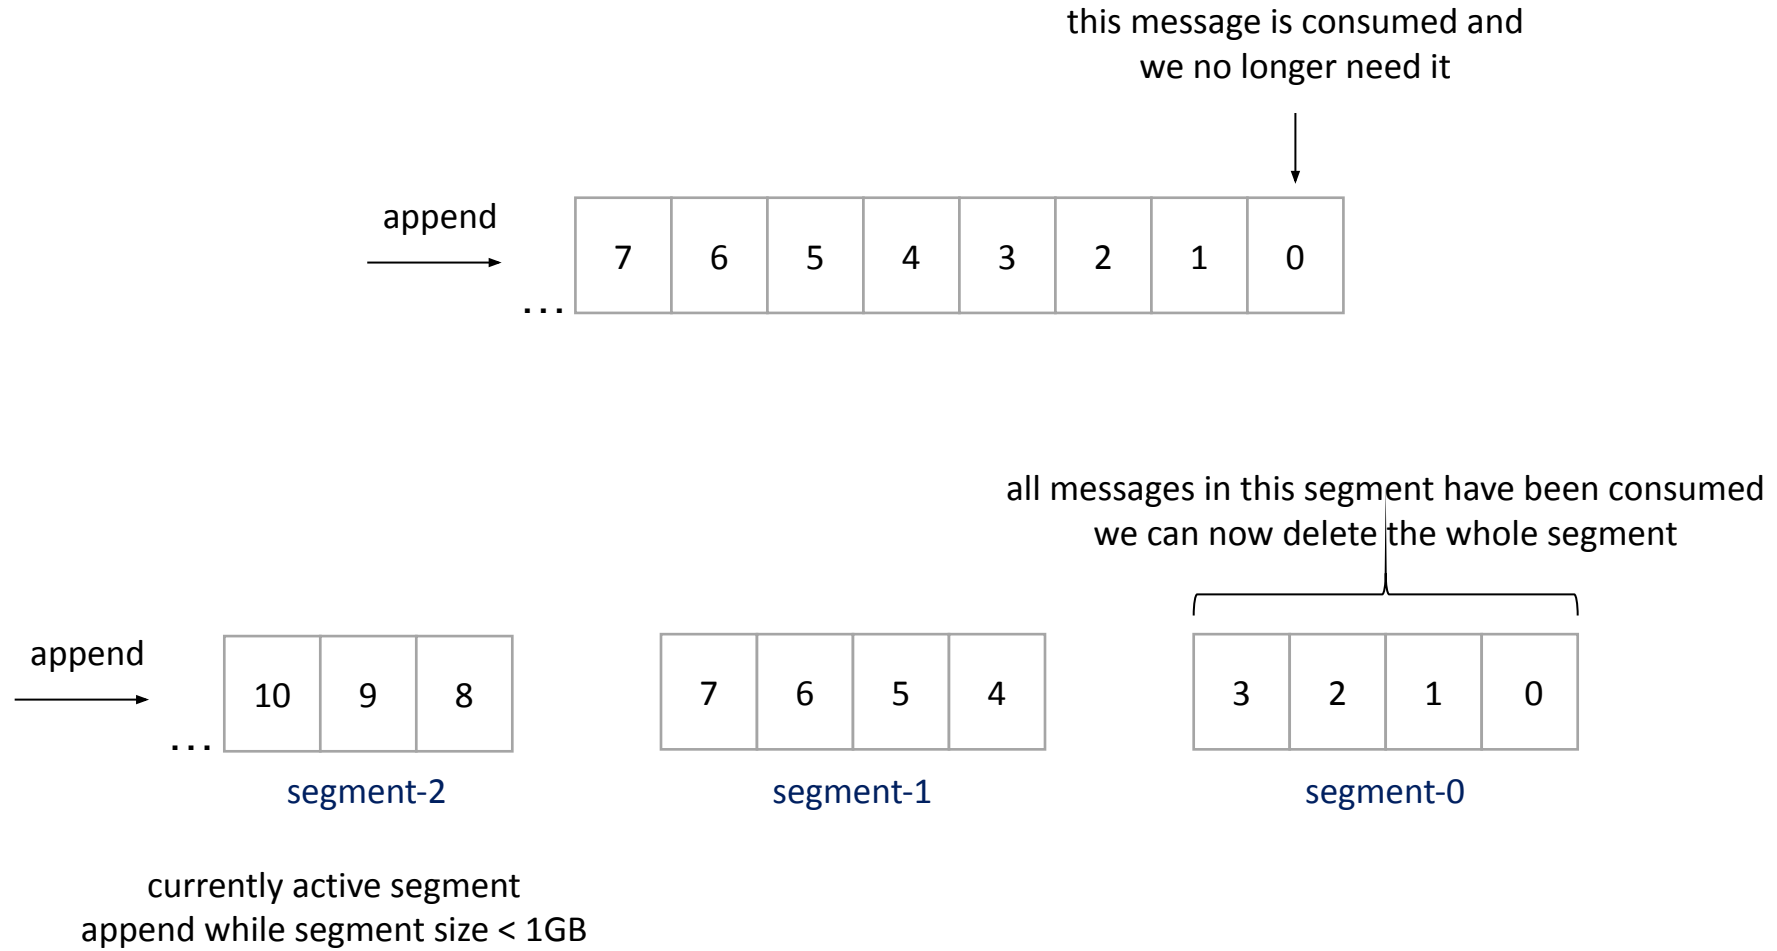


log


(data log, journal)

append-only and totally-ordered sequence of records

# Log



# Log



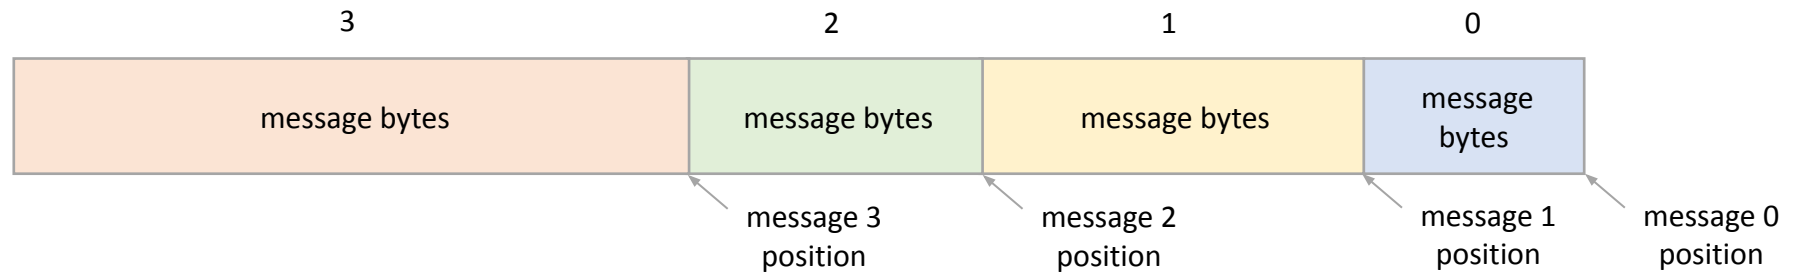
naïve approach

iterate over segments and messages starting with the oldest segment



better approach

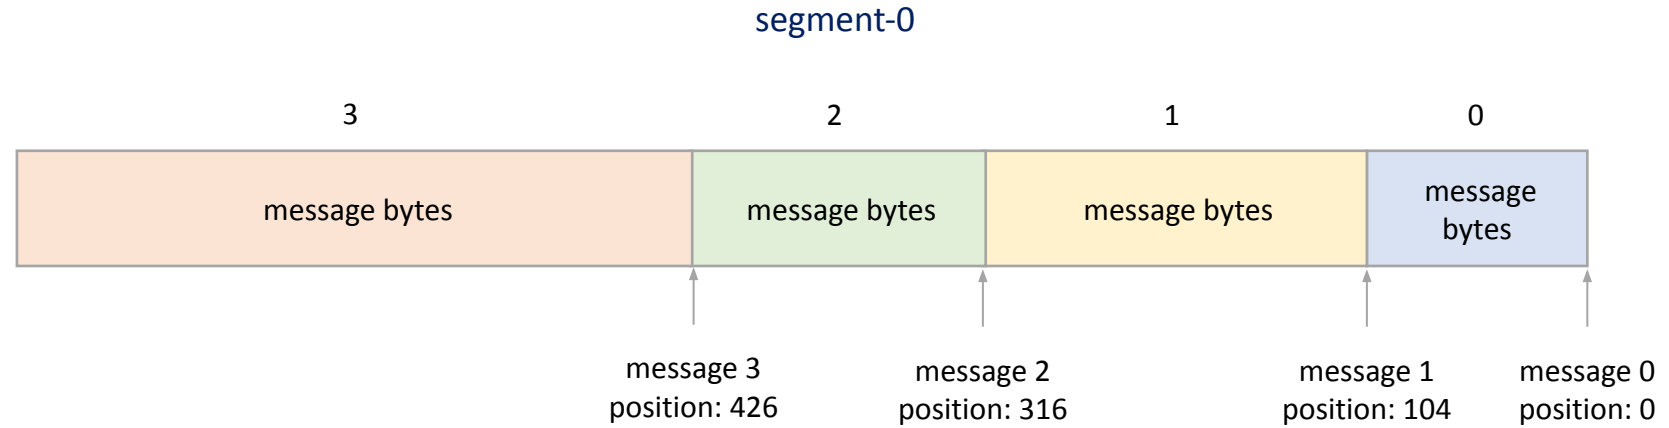
for each message know its position in the segment file



# Log

mapping

Message Id	Segment	Position
0	segment-0	0
1	segment-0	104
2	segment-0	316
3	segment-0	426
...	...	...



**where** and **how**  
do we store this mapping?