

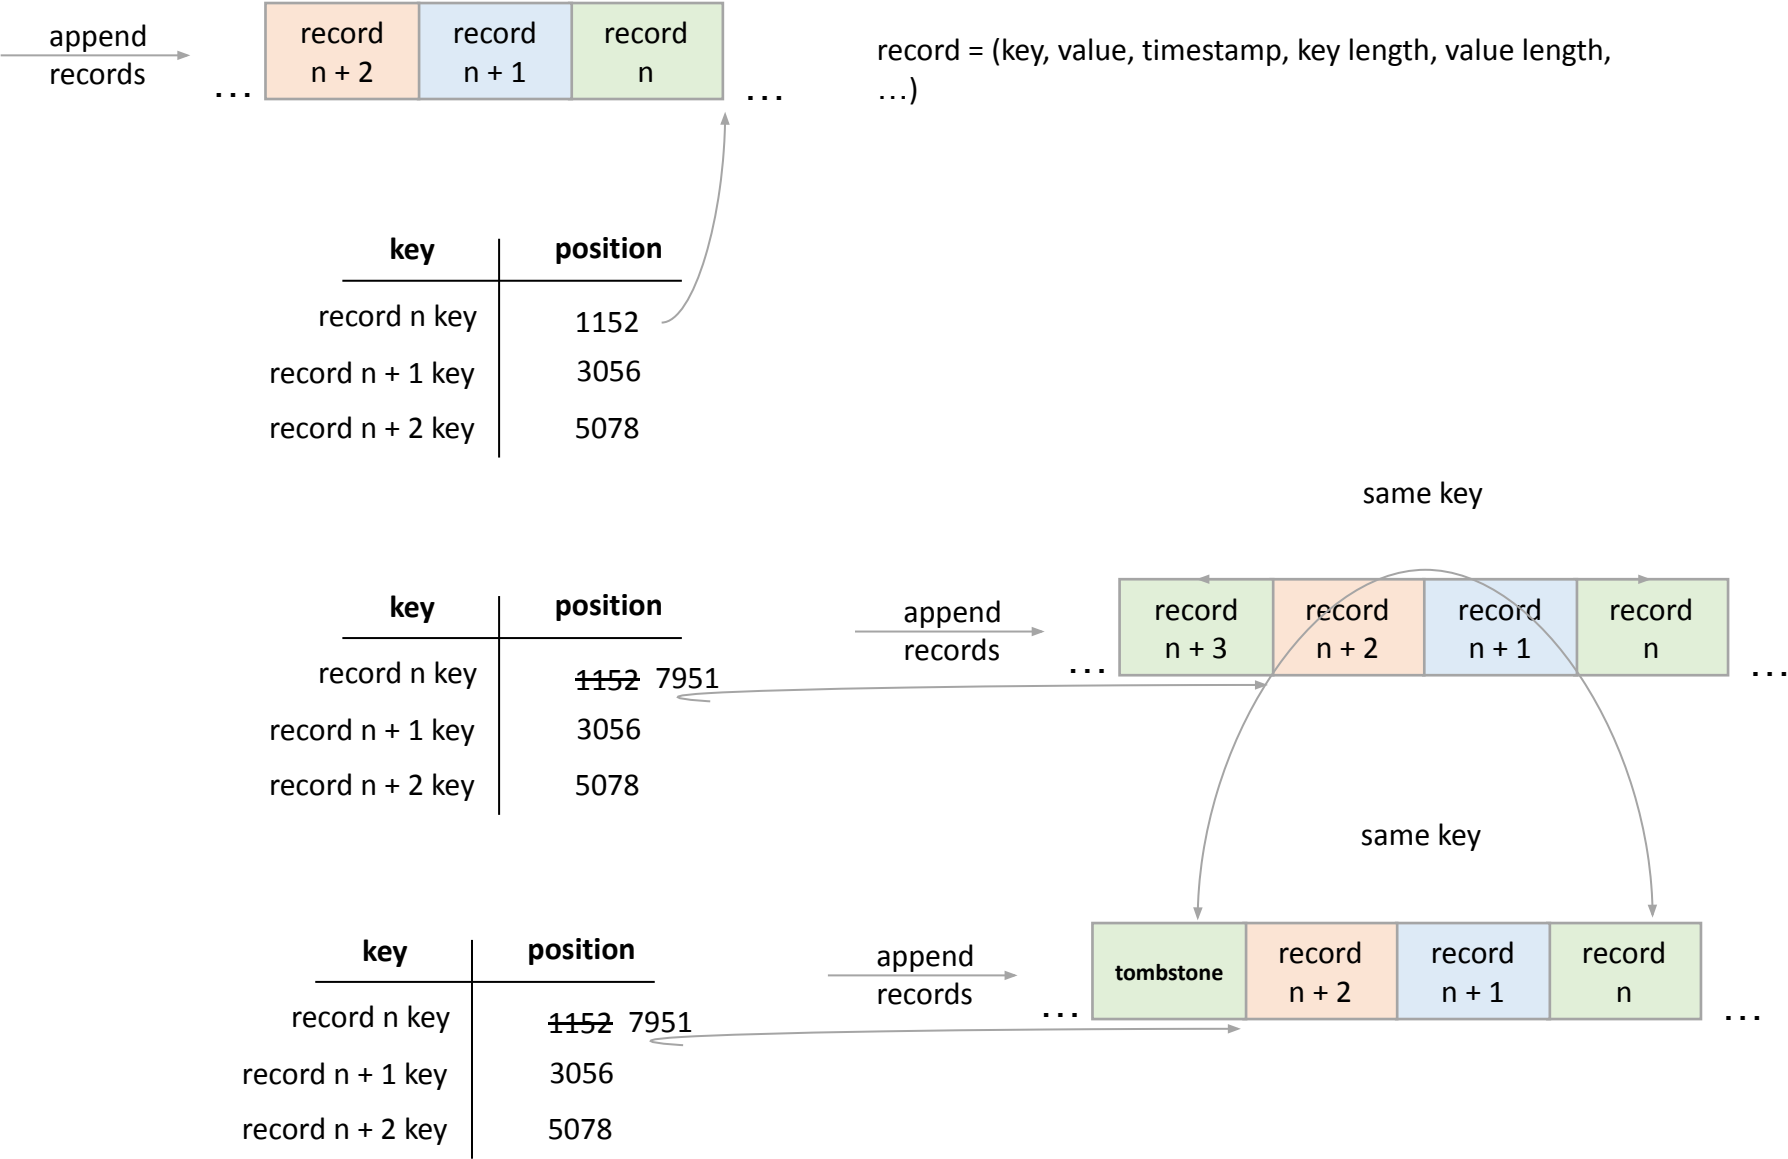
Simple key-value database

Create


Read

Update

Delete



Simple key-value database



Let me express some concerns about this design...



software engineer

1. records are constantly appended

Potential problem: not enough disk space.

2. we may not have enough memory to store the entire index

Potential problem: high query latency because indexes must be loaded from disk.

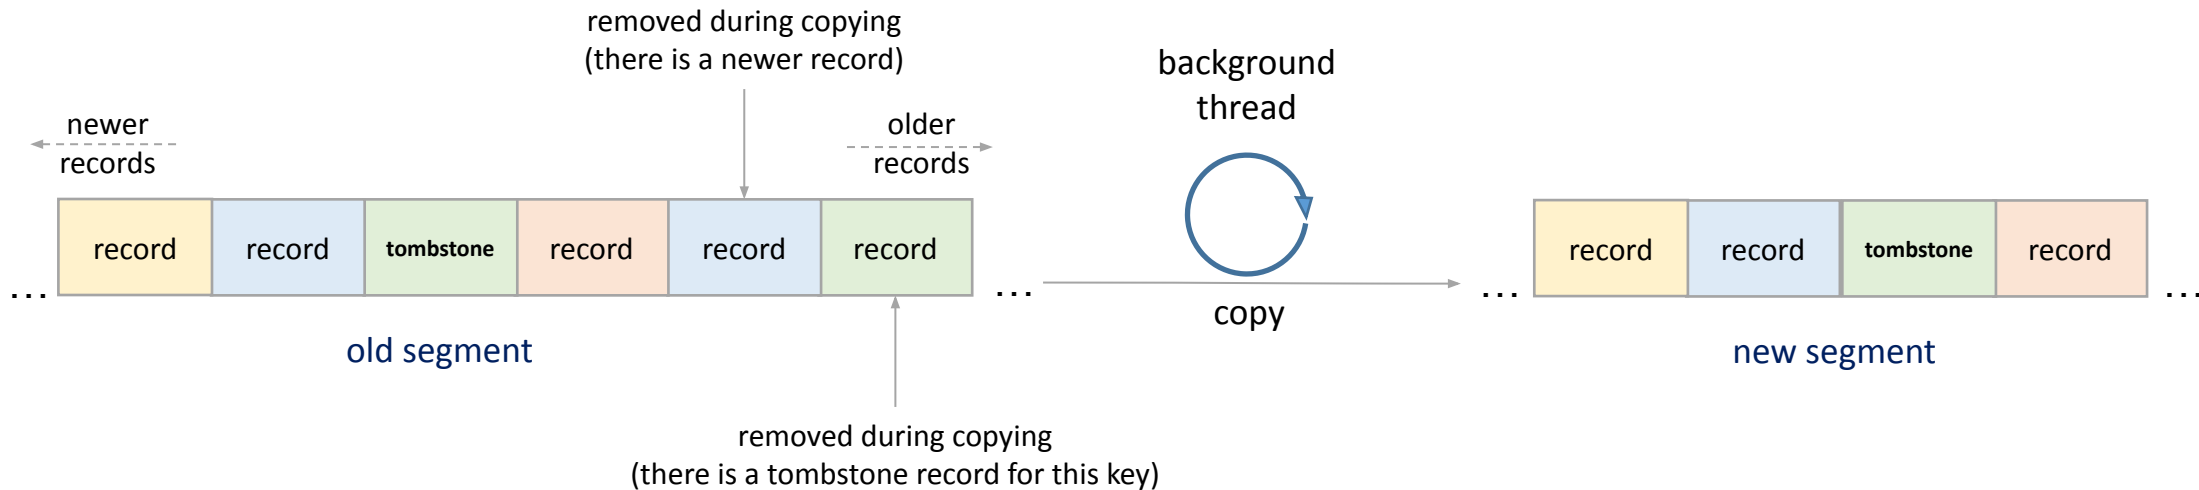
3. poor support for range queries

Simple key-value database

log compaction

the most recent record for a key is retained,
while older records for the same key are deleted

The default storage engine in Riak, a distributed NoSQL key-value data store, is based on ideas discussed in this video.



- old (bigger) segment is replaced by the new (smaller) segment
- several smaller segments can be merged together into one segment