

Batching

send multiple messages with a single request

benefits

- **increased throughput**
less overhead for HTTP headers, more efficient use of threads and connections
- **decreased costs**
the fewer requests we make, the less money we pay for using services with a pay-per-request pricing model
- **less chance of being throttled**
services often protect themselves from being abused by clients that send too many requests in a short period of time

Batching

batching introduces complexity on both the client and server side

accumulate messages in a buffer and
flush the buffer based on time or size

messages are processed one by one and
partial failures are possible



can be hard to implement and configure correctly

handling partial failures can be tricky

Batching

how the server handles batch requests

- treat the entire request as a single atomic unit
request succeeds only when all nested operations complete successfully
- treat each nested operation independently and
report back failures for each individual operation
service processing the request tries to make as much progress as possible

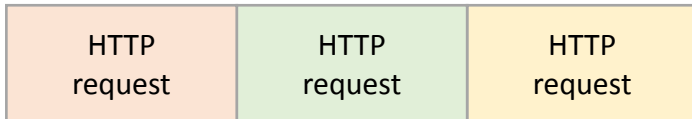
this approach is more common
in practice

Batching

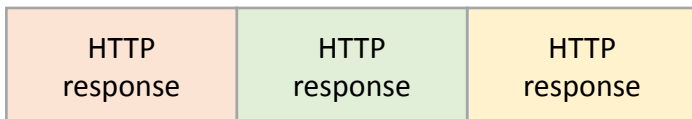
batch request format

set of n **requests** batched together

standard HTTP request



standard HTTP response



example

Google Drive batch API

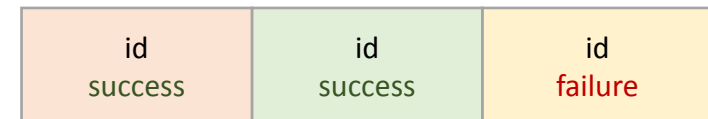
<https://developers.google.com/drive/api/v3/batch>

list of n **resources** batched together

standard HTTP request



standard HTTP response



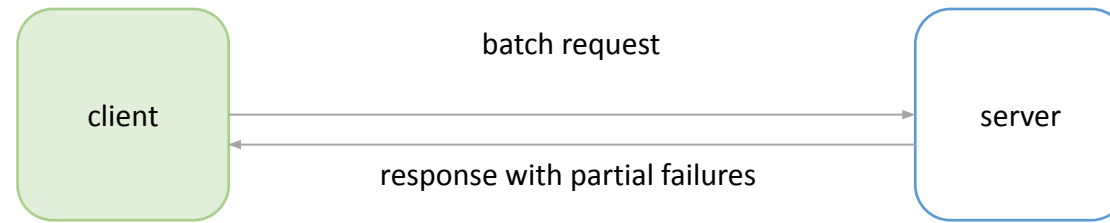
example

AWS SQS batch API

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/APIReference/API_SendMessageBatch.html

Batching

how the client handles failed nested operations



- retry the entire batch request
- retry each failed operation individually
- create another batch request containing only failed individual operations

