


# Partitioning strategies

lookup  
strategy

range  
strategy

hash  
strategy



How do we choose a shard  
for each user?

# Partitioning strategies

create and store a mapping between keys and shards

## lookup strategy

assign a shard randomly  
assign based on shards utilization

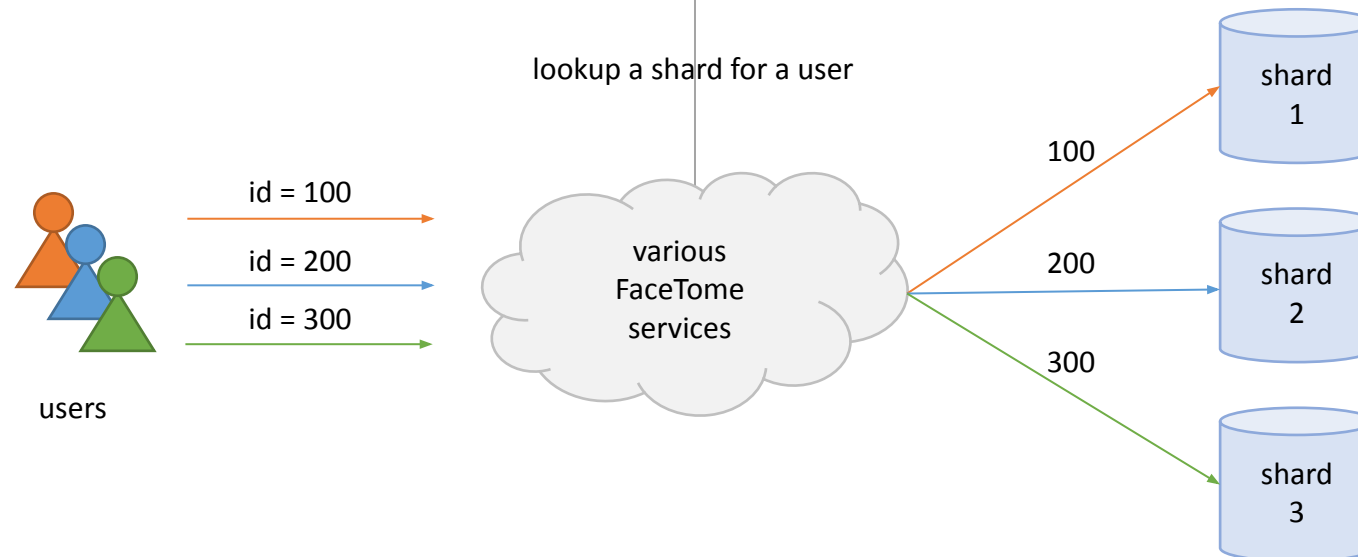
## lookup service

| User Id | Last Name | Shard Id |
|---------|-----------|----------|
| 100     | Apple     | shard 1  |
| 200     | Apricot   | shard 2  |
| 300     | Banana    | shard 3  |

How to make the lookup service

- highly available?
- fast?
- scalable?
- consistent?

- offers a lot of control over shard assignment logic
- hard dependency on the mapping
- requires highly-available and fast lookup service
- mapping can get really big over time (and needs to be partitioned)




# Partitioning strategies

each shard is responsible for a continuous range of keys

## range strategy

- easy to implement
- works well with range queries



- provides suboptimal balancing between shards (which may lead to the hot shard problem)

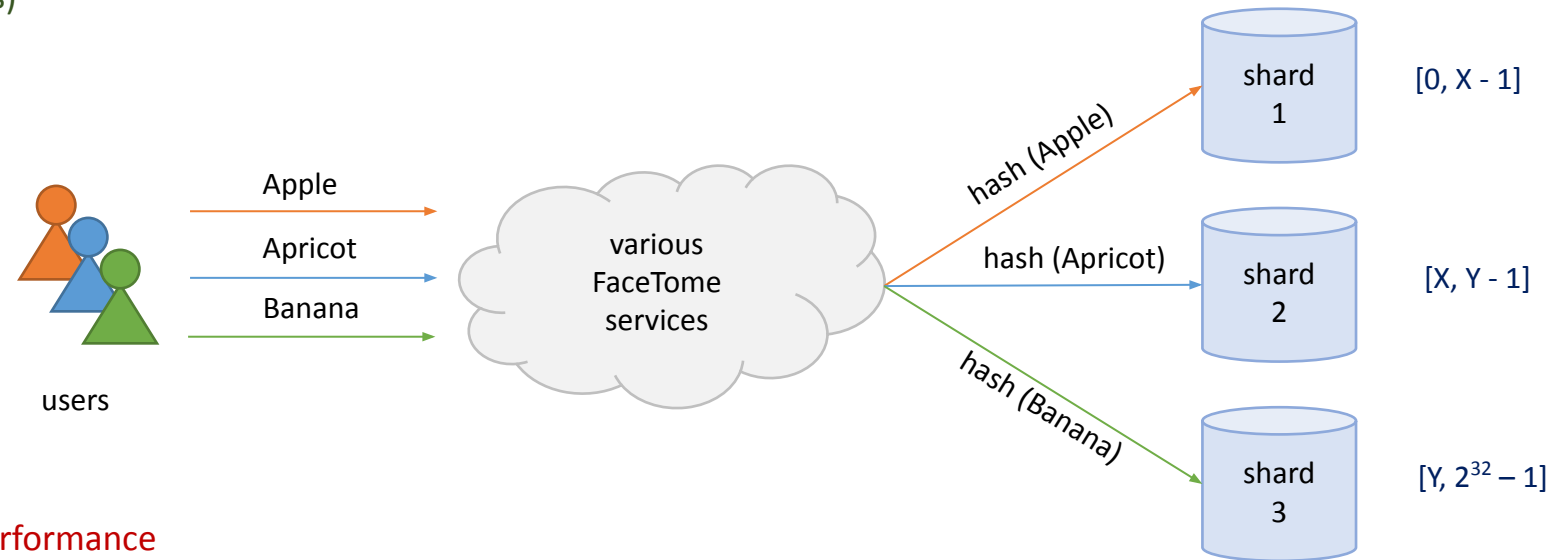
# Partitioning strategies

each shard is responsible for a continuous range of hashes

## hash strategy

hash (user last name)  $\rightarrow [0, 2^{32} - 1]$

- provides more even data distribution across shards  
(reduces the chance of hot spots)



- range queries may have poor performance
- computing the hash imposes an additional overhead

define shard boundaries by

- spacing them evenly
- using consistent hashing