Consistent hashing

to implement hash partitioning we need to define…

- boundaries for each shard (a range of hashes each shard owns)
- where each shard lives (which physical server each shard lives on)
- how to rebalance shards (pick and implement a rebalancing strategy)

consistent hashing

Consistent hashing

hash functions

hash function (key) ⟶ *integer* $[0, 2^{31})$

1. How to split this range into smaller hash value intervals (shards).
2. How to assign each interval (shard) to a server.

cryptographic hash functions
(e.g. SHA)

non-cryptographic hash functions
(e.g. MurmurHash)

Consistent hashing

consistent hashing ring

server 1

server 2

server 3

server 4

server 5

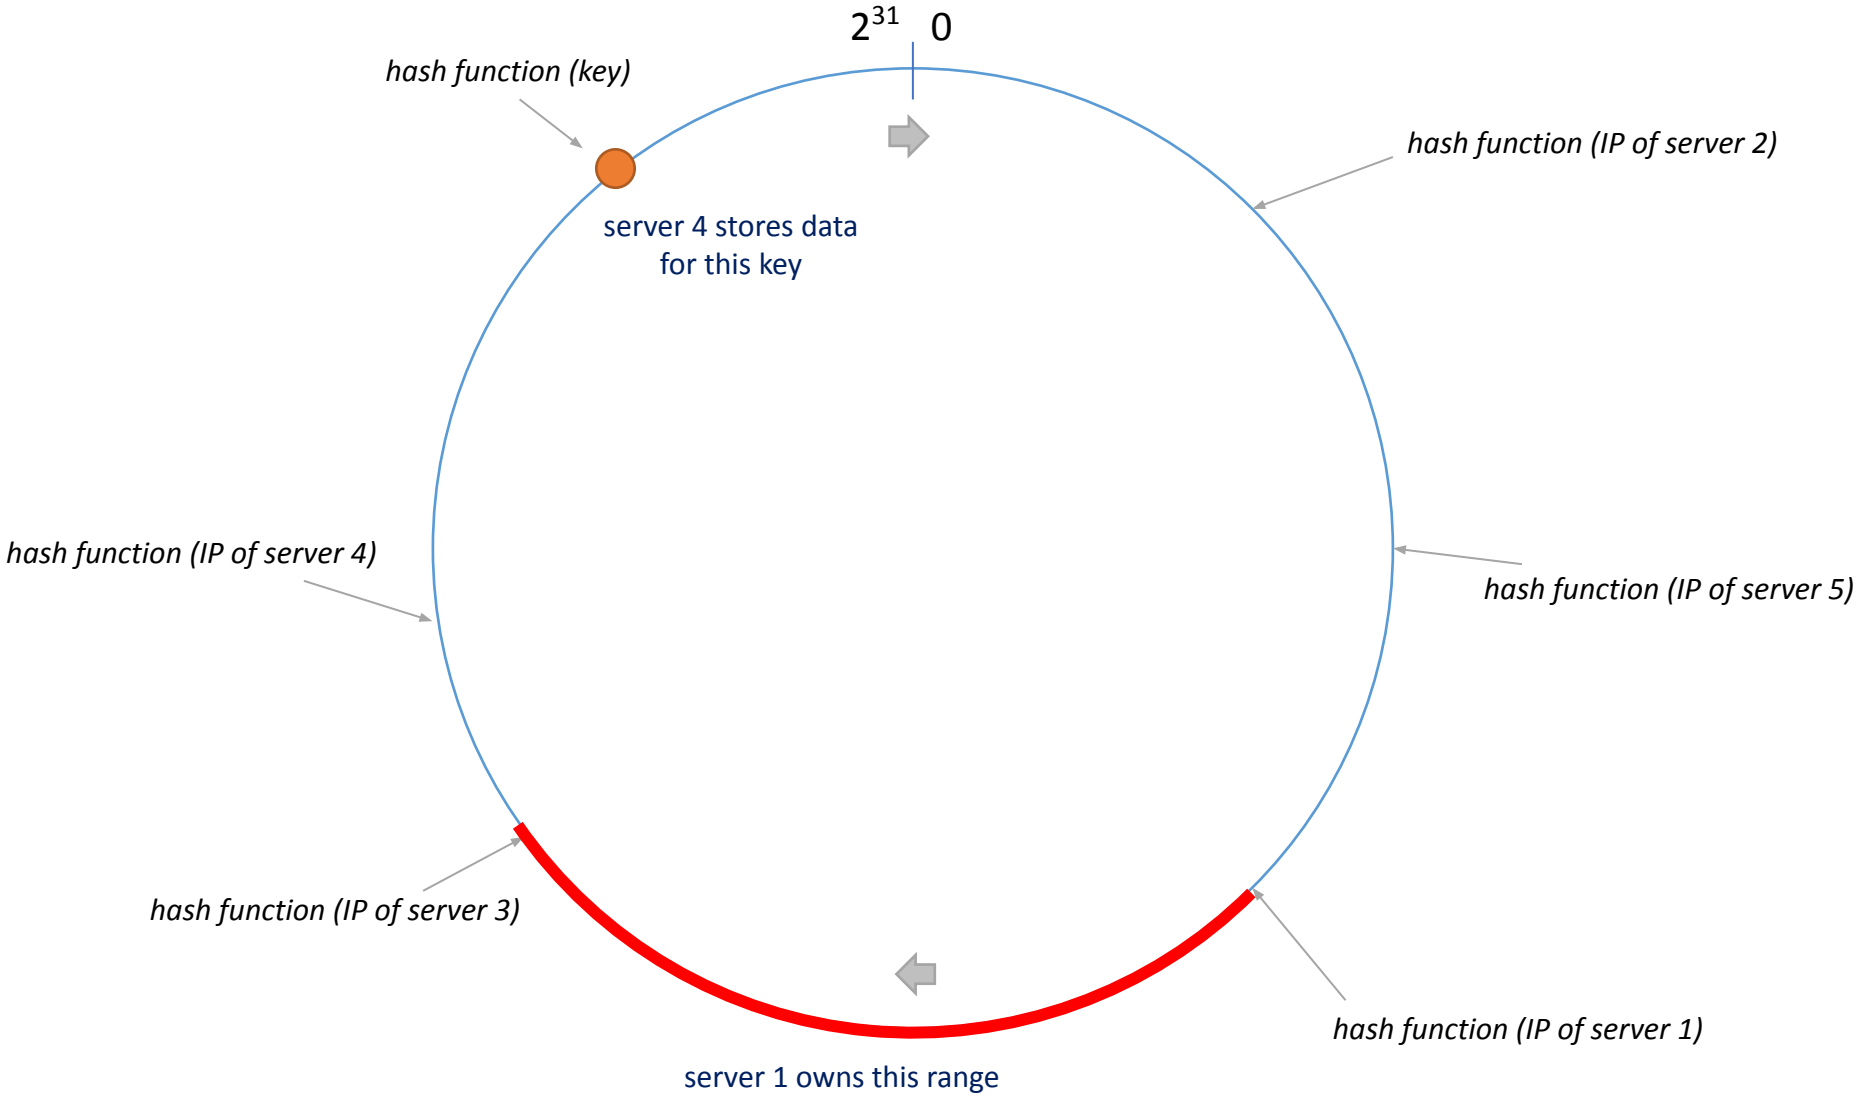

Consistent hashing

implementation details

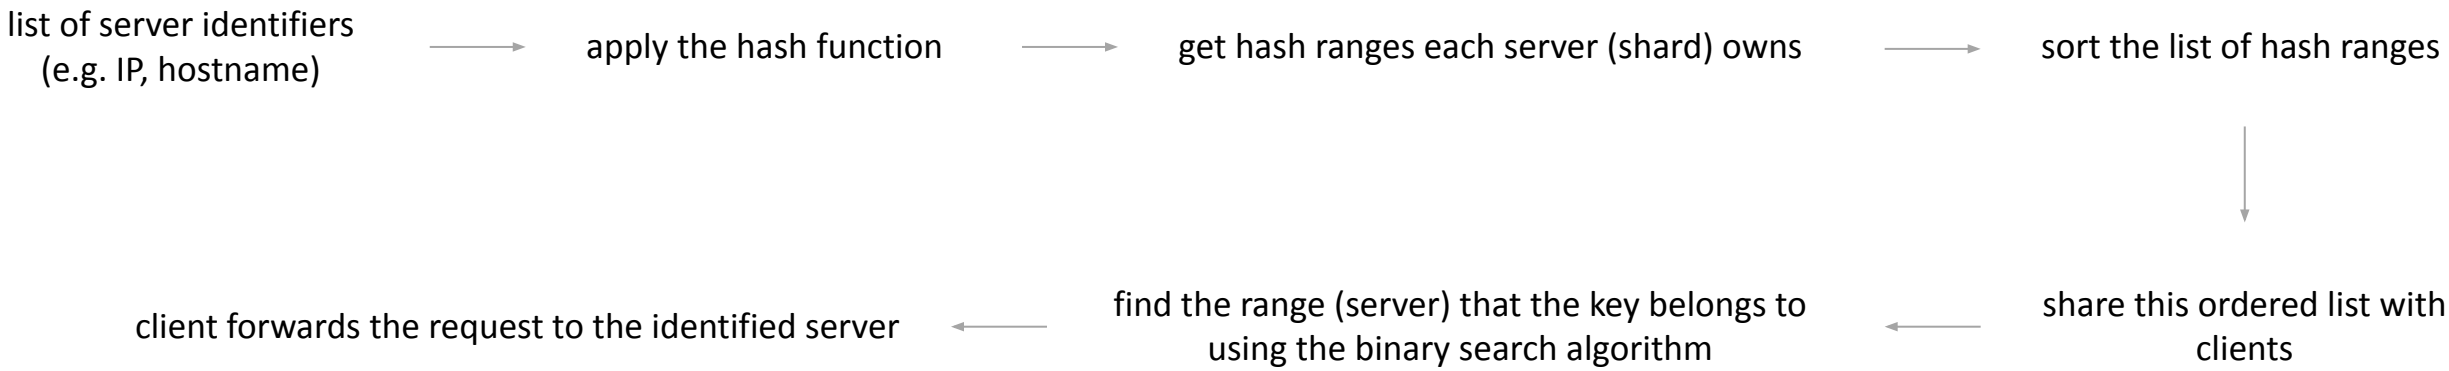

Consistent hashing

rebalancing

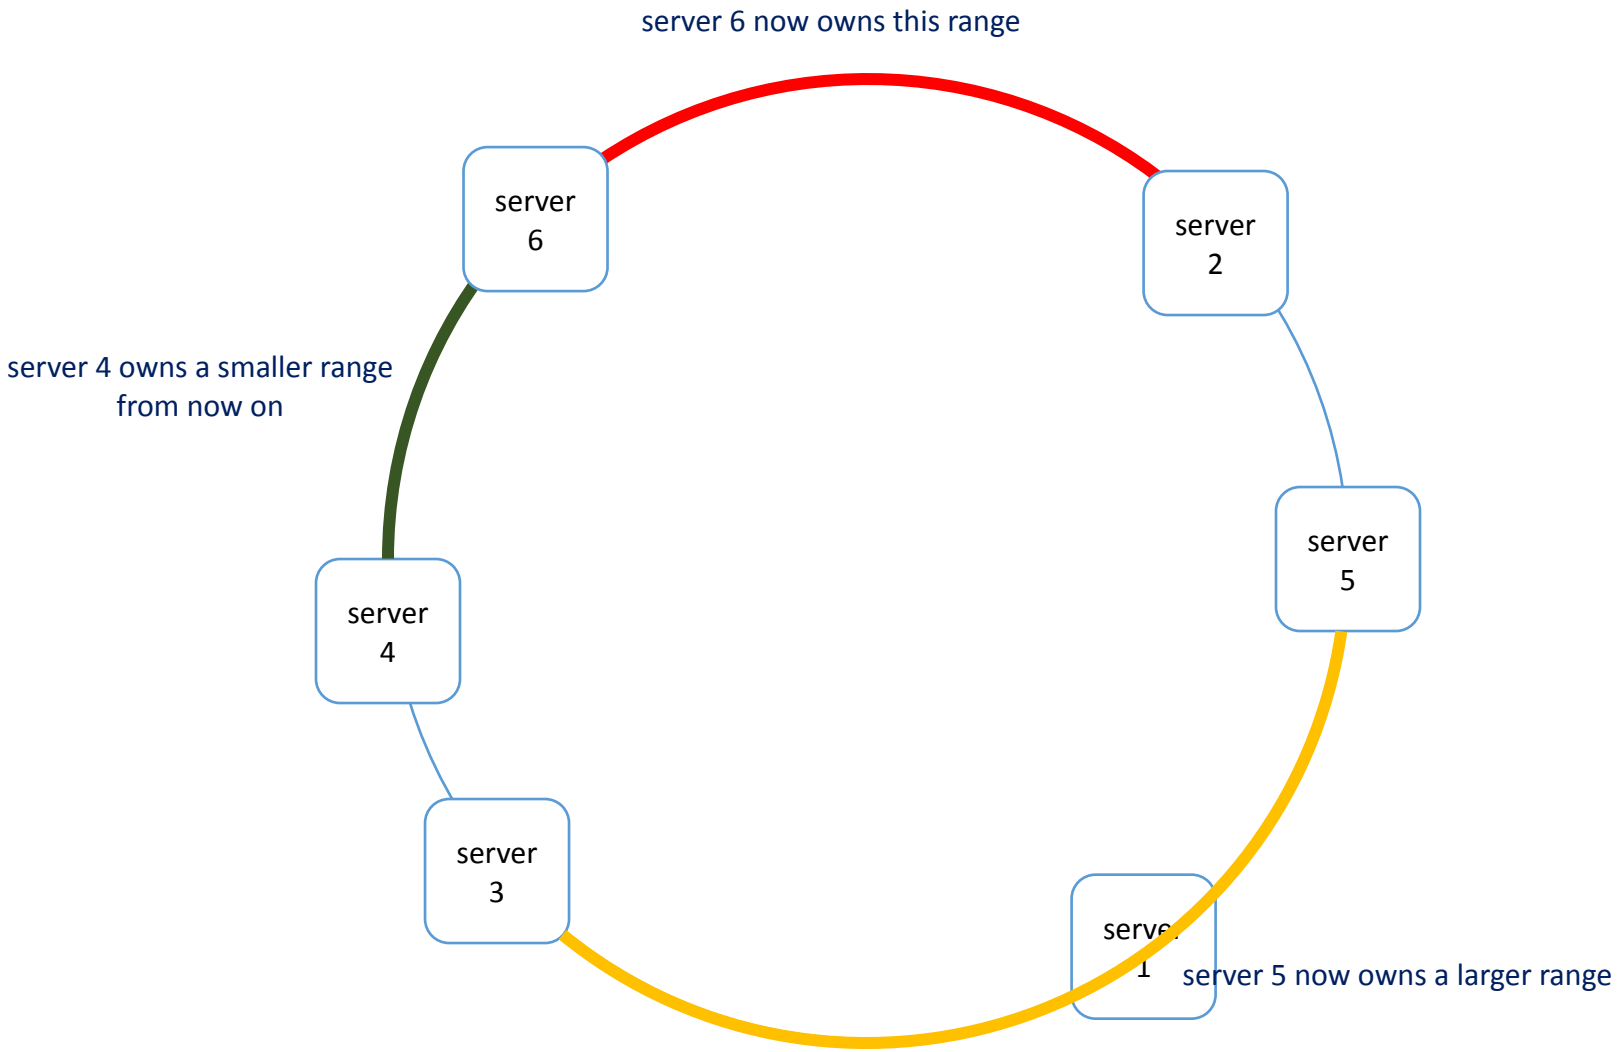

Consistent hashing

fixed number of shards

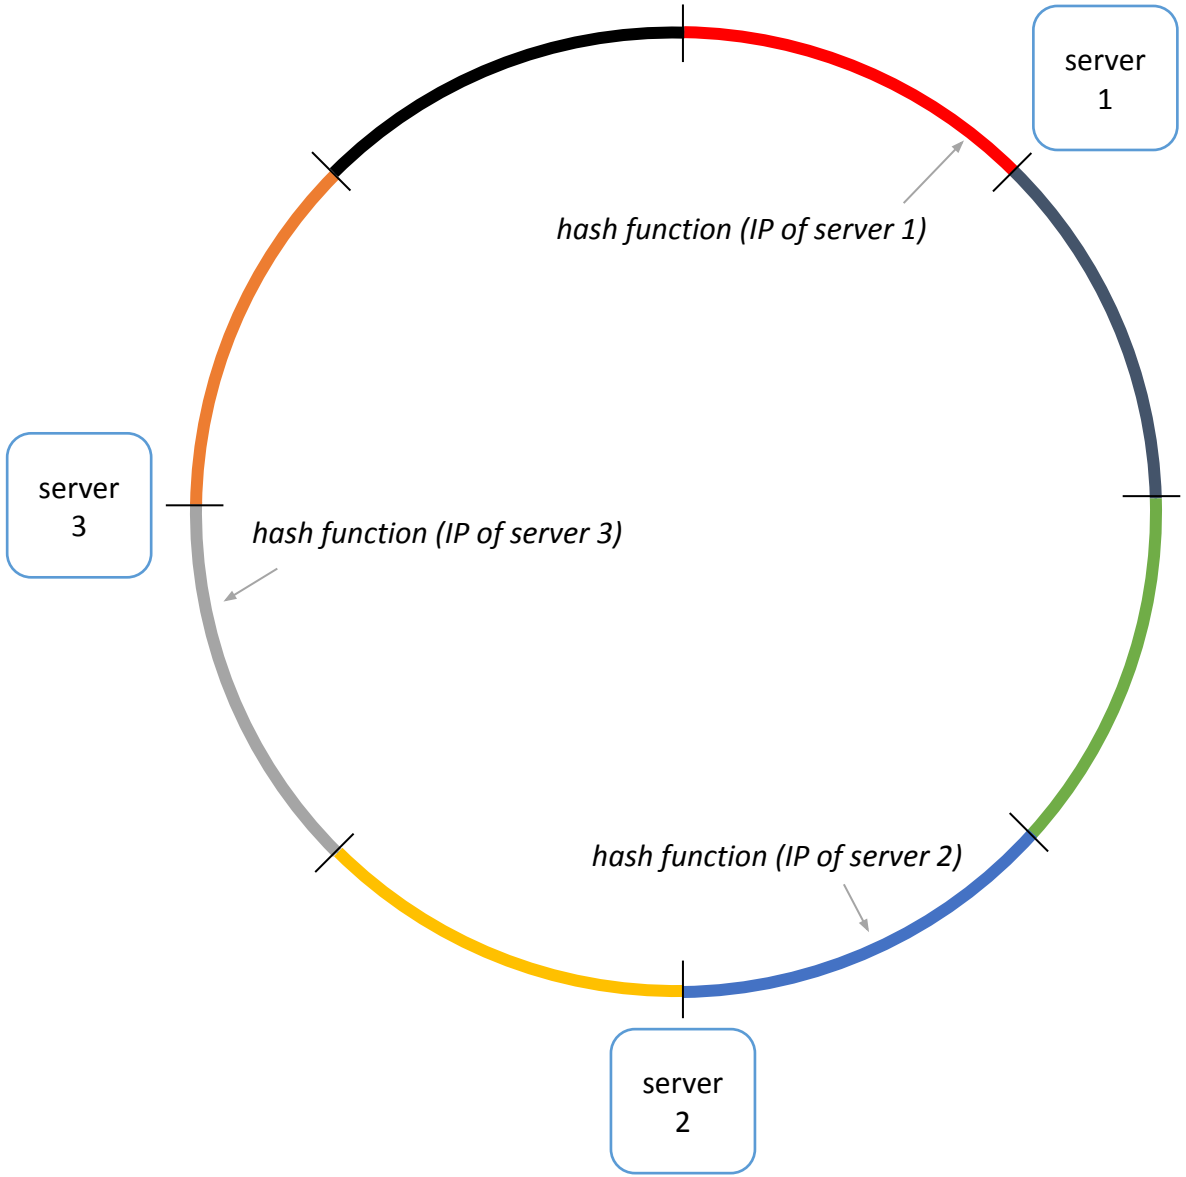

Consistent hashing

shard replicas

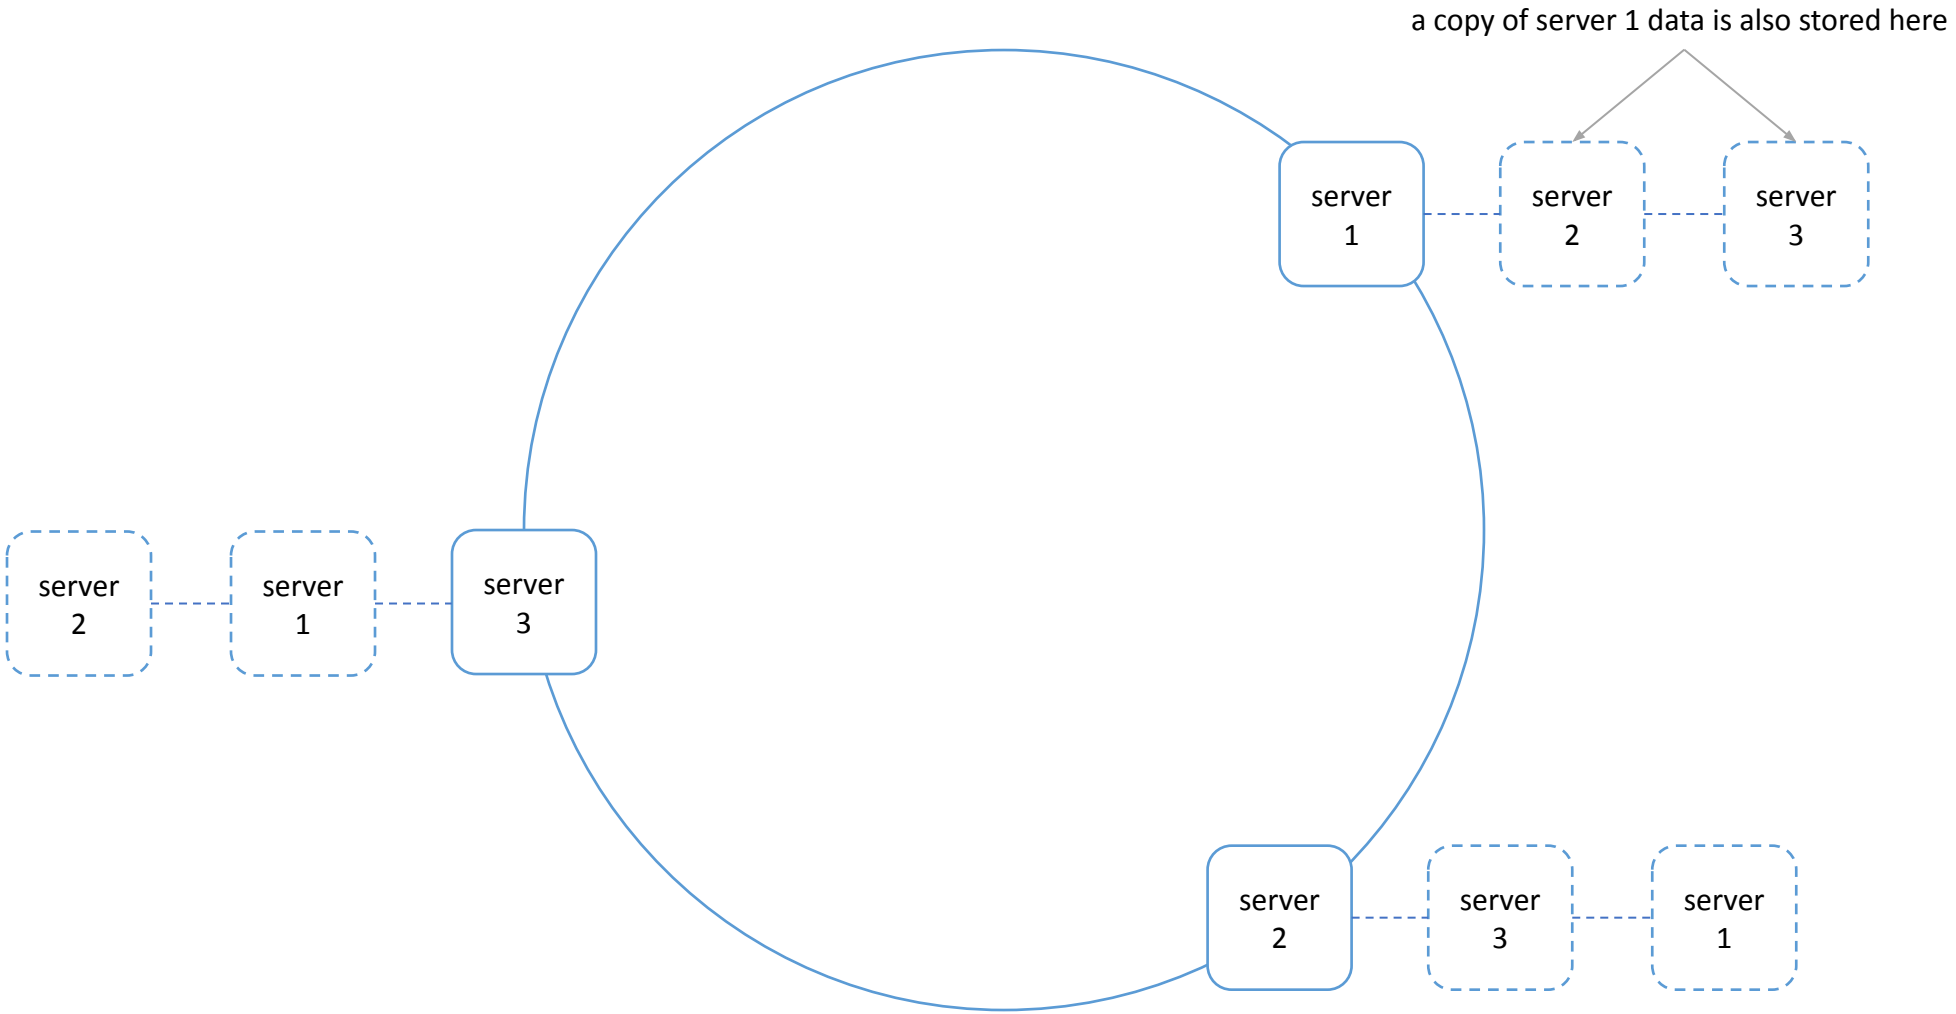

Consistent hashing

disadvantages

- domino effect
- servers do not split the circle evenly

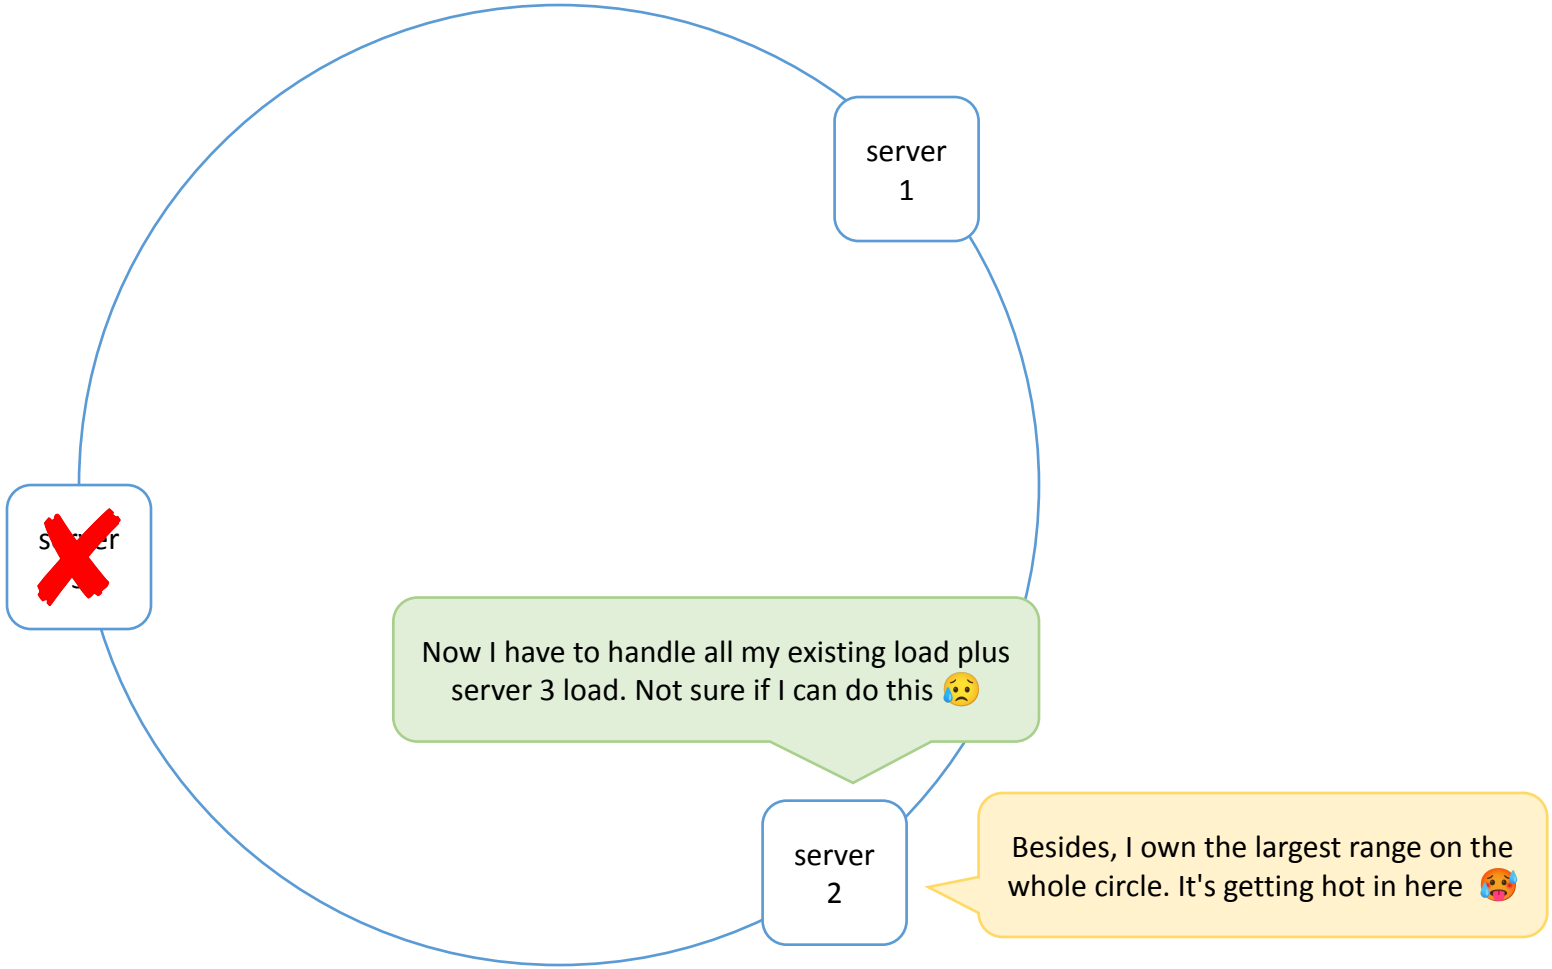

Consistent hashing

virtual nodes

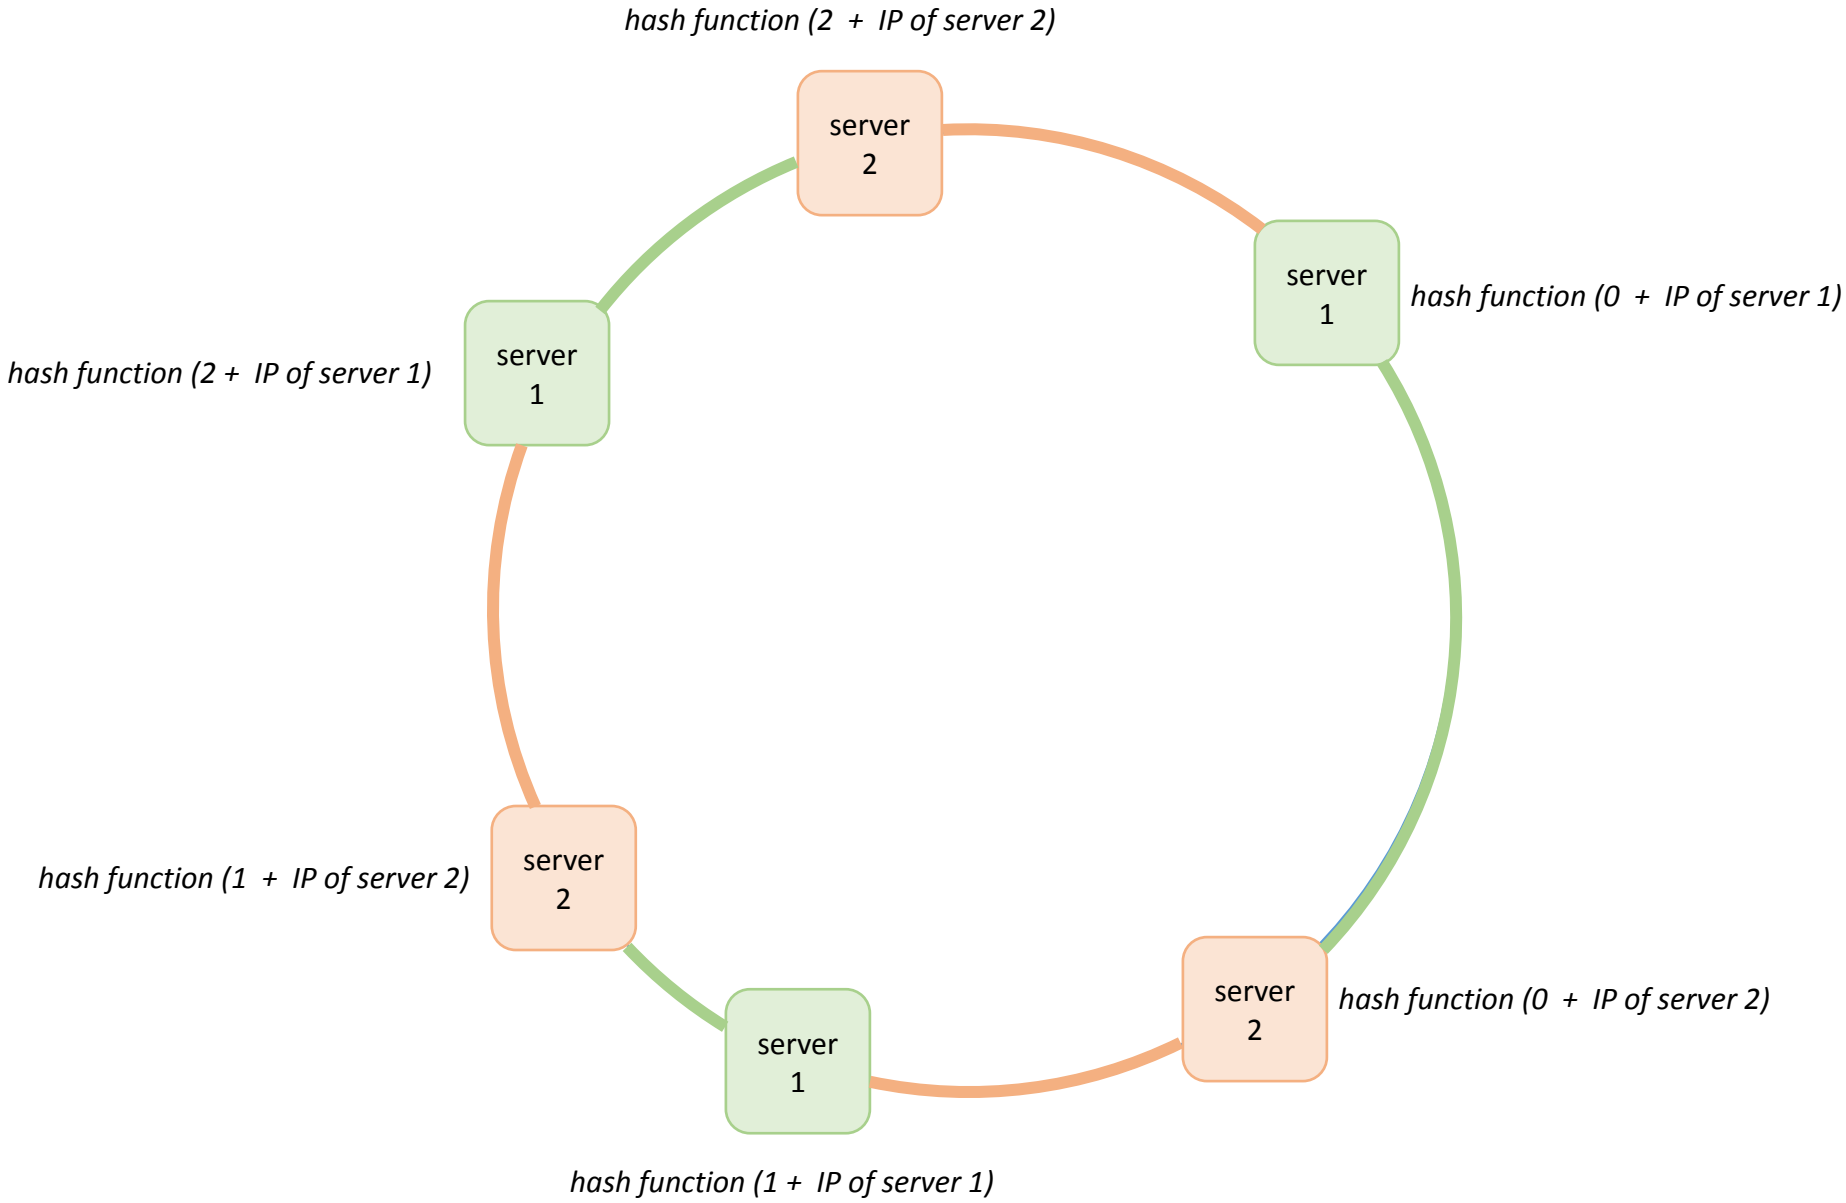

Consistent hashing

virtual nodes

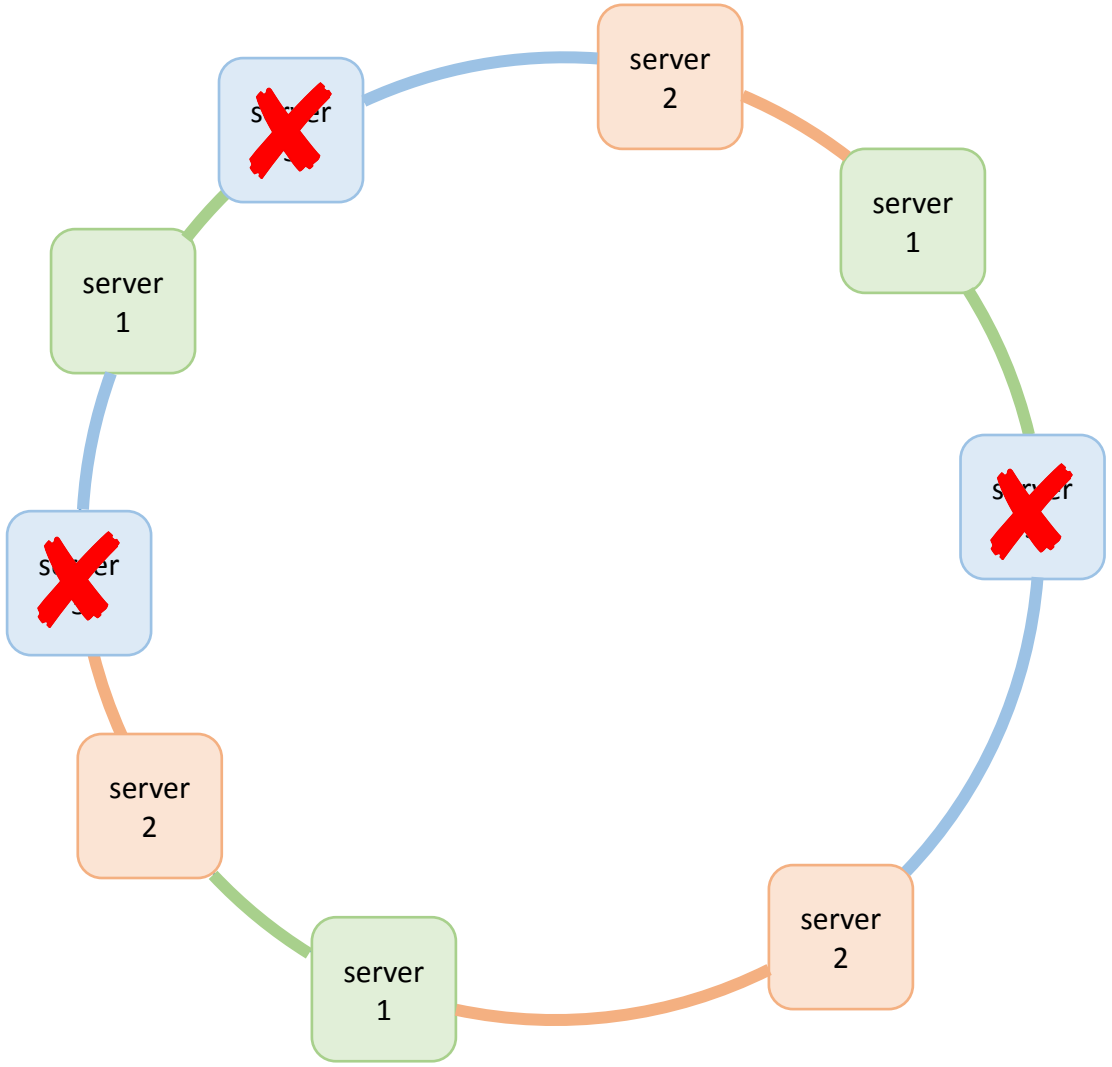

Consistent hashing

examples

- databases (Cassandra, Couchbase, Riak, Voldemort)
- distributed caches (client libraries, e.g. Ketama)
- content delivery networks (Akamai)
- network load balancers (Maglev)
- chat applications (Discord)
- messaging systems (RabbitMQ)