

[Log In](#)

[Join](#)

[Back To Course Home](#)

Grokking Modern System Design Interview for Engineers & Managers

0% completed

System Design Interviews

Introduction

Abstractions

Non-functional System Characteristics

Back-of-the-envelope Calculations

Building Blocks

Domain Name System

Load Balancers

Databases

- ① Introduction to Databases
- Types of Databases
- Data Replication
- Data Partitioning

Trade-offs in Databases

Key-value Store

Content Delivery Network (CDN)

Sequencer

Distributed Monitoring

Monitor Server-side Errors

Monitor Client-side Errors

Distributed Cache

Distributed Messaging Queue

Pub-sub

Rate Limiter

Blob Store

Distributed Search

Distributed Logging

Distributed Task Scheduler

Sharded Counters

Concluding the Building Blocks Discussion

Design YouTube

Design Quora

Design Google Maps

Design a Proximity Service / Yelp

Design Uber

Design Twitter

Design Newsfeed System

Design Instagram

Design a URL Shortening Service / TinyURL

Design a Web Crawler

Design WhatsApp

Design Typeahead Suggestion

Design a Collaborative Document Editing Service / Google Docs

Spectacular Failures

Concluding Remarks

Course Certificate

Mark Course as Completed

Trade-offs in Databases

Learn when to use horizontal sharding instead of vertical sharding and vice versa.

We'll cover the following

- Which is the best database sharding approach?
 - Advantages and disadvantages of a centralized database
 - Advantages
 - Disadvantages
 - Advantages and disadvantages of a distributed database
 - Advantages
 - Disadvantages
- Query optimization and processing speed in a distributed database
 - Parameters assumption
 - Possible approaches
- Conclusion

Which is the best database sharding approach?#

Both horizontal and vertical sharding involve adding resources to our computing infrastructure. Our business stakeholders must decide which is suitable for our organization. We must scale our resources accordingly for our organization and business to grow, to prevent downtime, and to reduce latency. We can scale these resources through a combination of adjustments to CPU, physical memory requirements, hard disk adjustments, and network bandwidth.

The following sections explain the pros and cons of no sharding versus sharding.

Advantages and disadvantages of a centralized database#

Advantages#

- Data maintenance, such as updating and taking backups of a centralized database, is easy.
- Centralized databases provide stronger consistency and ACID transactions than distributed databases.
- Centralized databases provide a much simpler programming model for the end programmers as compared to distributed databases.
- It's more efficient for businesses to have a small amount of data to store that can reside on a single node.

Disadvantages#

- A centralized database can slow down, causing high latency for end users, when the number of queries per second accessing the centralized database is approaching single-node limits.
- A centralized database has a single point of failure. Because of this, its probability of not being accessible is much higher.

Advantages and disadvantages of a distributed database#

Advantages#

- It's fast and easy to access data in a distributed database because data is retrieved from the nearest database shard or the one frequently used.
- Data with different **levels of distribution transparency** can be stored in separate places.

Intensive transactions consisting of queries can be divided into multiple optimized subqueries, which can be processed in a parallel fashion.

Disadvantages#

- Sometimes, data is required from multiple sites, which takes more time than expected.
- Relations are partitioned vertically or horizontally among different nodes. Therefore, operations such as joins need to reconstruct complete relations by carefully fetching data. These operations can become much more expensive and complex.
- It's difficult to maintain consistency of data across sites in the distributed database, and it requires extra measures.
- Updates and backups in distributed databases take time to synchronize data.

Query optimization and processing speed in a distributed database#

A transaction in the distributed database depends on the type of query, number of sites (shards) involved, communication speed, and other factors, such as underline hardware and the type of database used. However, as an example, let's assume a query accessing three tables, **Store**, **Product**, and **Sales**, residing on different sites.

The number of attributes in each table is given in the following figure:

Database schema consisting of three tables: Store, Product, and Sales

Let's assume the distribution of both tables on different sites is the following:

- The **Store** table has 10,000 tuples stored at site A.
- The **Product** table has 100,000 tuples stored at site B.
- The **Sales** table has one million tuples stored at site A.

Now, assume that we need to process the following query:

```
Select Store_key from (Store JOIN Sales JOIN Product)
where Region= 'East' AND Brand='Wolf';
```

The above query performs the join operations on the **Store**, **Sales**, and **Product** tables and retrieves the **Store_key** values from the table generated in the result of join operations.

Next, assume every stored tuple is 200 bits long. That's equal to 25 Bytes. Furthermore, estimated cardinalities of certain intermediate results are as follows:

- The number of the **Wolf** brand is 10.
- The number of **East** region stores is 100,000.

Communication assumptions are the following:

- Data rate = 50M bits per second
- Access delay = 0.1 second

Parameters assumption#

Before processing the query using different approaches, let's define some parameters:

a = Total access delay

b = Data rate

v = Total data volume

Now, let's compute the total communication time, T , according to the following formula:

$$T = a + \frac{v}{b}$$

Let's try the following possible approaches to execute the query.

Possible approaches#

- Move the **Product** table to site A and process the query at A.

$$T = 0.1 + \frac{100,000 \times 200}{50,000,000} = 0.5 \text{ seconds}$$

Here, 0.1 is the access delay of the table on site A, and 100,000 is the number of tuples in the **Product** table. The size of each tuple in bits is 200, and 50,000,000 is the data rate. The 200 and 50,000,000 figures are the same for all of the following calculations.

- Move **Store** and **Sales** to site B and process the query at B:

$$T = 0.2 + \frac{(10,000 + 1,000,000) \times 200}{50,000,000} = 4.24 \text{ seconds}$$

Here, 0.2 is the access delay of the **Store** and **Product** tables. The numbers 10,000 and 1,000,000 are the number of tuples in the **Store** and **Product** tables, respectively.

- Restrict **Brand** at site B to **Wolf** (called projection) and move the result to site A:

$$T = 0.1 + \frac{10 \times 200}{50,000,000} \approx 0.1 \text{ seconds}$$

Here, 0.1 is the access delay of the **Product** table. The number of the **Wolf** brand is 10, hence the number of tuples.

When we compare the three approaches, the third approach provides us the least latency (0.1 seconds). This example shows that careful query optimization is also critical in the distributed database.

Conclusion#

Data distribution (vertical and horizontal sharding) across multiple nodes aims to improve the following features, considering that the queries are optimized:

- Reliability (fault-tolerance)
- Performance
- Balanced storage capacity and dollar costs

Both centralized and distributed databases have their pros and cons. We should choose them according to the needs of our application.

Back

Data Partitioning

Next

System Design: The Key-value Store

Mark as Completed

Report an Issue